

Erode AI Course Syllabus

Selvakumar Murugan

December 19, 2025

Contents

Syllabus - AI / LLM Engineering with Transparent Abstractions

Summary

I am planning the course in such way that it trains students to **build, evaluate, and deploy reliable AI/LLM systems without myth, hype, or fragile intuitions**.

We don't want to focus on:

- prompt tricks,
- surface-level tooling,
- or abstract ML math disconnected from practice,

Instead we stress:

- **correct mental models,**
- **engineering discipline,**

My intention is to make students understand what these systems are, what they are not, where they fail, and how to control them. not do not merely **use** LLMs. But we still design the course around LLMs, and cover ML concepts spread over the entire syllabus.

- This is not a**
- an AI hype course
 - a prompt engineering bootcamp

- a math-heavy ML theory course

But • a rigorous, honest, industry-aligned program
 • to build and explain AI systems.

The course covers • probabilistic modeling

- training vs inference
- loss functions
- optimization
- generalization
- evaluation

But always framed as, **machinery behind behavior**, not explanations of intelligence. Students can later learn deeper math **without unlearning incorrect metaphors**.

Goals

Most AI/LLM courses suffer from one or more of these failures:

- Anthropomorphic explanations (“the model understands”)
- Overconfidence in prompts and benchmarks
- Confusion between training and inference
- Ignoring uncertainty, failure modes, and deployment realities
- Teaching ML math that students later misuse conceptually
- Producing candidates who sound impressive but are unsafe in production

Industry consequence: Teams spend months unlearning bad intuitions, fixing fragile systems, and correcting overclaims made by well-meaning but undertrained engineers.

I’d like our course be designed explicitly to prevent that.

Core Principles

- Mental Model Hygiene Students must be rigorously trained to:
 - avoid anthropomorphism
 - distinguish model behavior from system behavior
 - treat uncertainty as inherent, not as error
 - use precise, defensible language

- Three Layer Abstractions that are Transparent Every concept is taught at three explicit levels:

1. **Behavioral** — what it appears to do
2. **Mechanical** — what actually happens
3. **Formal / Mathematical** — optional depth, clearly scoped

Students should always grasp the following to prevent long term confusion and re-learning in future:

- where the abstraction boundary is
- what is literal vs metaphorical, e.g: words like meaning, memory and reasoning are abused
- what machinery sits underneath e.g: Neural networks and Transformers

- Metaphor Audits The course actively identifies and dismantles misleading metaphors such as:

- "the model understands"
- "the agent plans"
- "the AI remembers"
- "RAG gives knowledge"

Students should learn to after going through the course:

- detect bad metaphors
- explain why they are dangerous
- replace them with precise language

This is a **very rare but critical professional skill even among most experienced engineers.**

- Big picture view instead of erratic metaphors that needs to unlearned
Students are trained to reason about:
 - pipelines
 - data flow
 - state
 - boundaries
 - failure propagation
 - glue code

LLMs are treated correctly as, probabilistic components inside engineered systems and not as intelligent entities.

- Just-In-Time Foundations (No Prerequisites) Students can enter with:
 - no programming background
 - no ML background
 - no math background

Foundations (Python, statistics, ML, information theory) are introduced so while avoiding intimidation:

- **only when needed**
- **only to the depth required**
- **with explicit warnings against misuse**

What the Syllabus Covers (High-Level)

Weeks 1–3: Core Reality of LLMs

- What LLMs actually are (probabilistic text generators)
- Prompting as conditioning, not programming
- Verification, tools, and uncertainty

Week 4: Programming Fundamentals (AI-Assisted)

- Code vs execution
- Determinism vs probability
- Control flow, functions, and verification

Week 5: APIs and Inference

- Stateless inference
- Cost, latency, and failure
- Secrets and configuration

Week 6: Guardrails and Structured Outputs

- Validation, retries, and contracts
- Why structure is enforced externally

Week 7: Retrieval-Augmented Generation (RAG)

- Retrieval vs understanding
- Embeddings as geometry
- Why RAG can help or harm

Week 8: Evaluation and Measurement

- Metrics as estimates
- Variance, sampling, and uncertainty
- Why benchmarks lie

Weeks 8.5–10.5: Traditional ML Foundations (Optional but Integrated)

- Learning as parameter optimization
- Loss functions and gradient descent (machinery, not mysticism)
- Overfitting, generalization, and evaluation
- Explicit mapping to standard ML theory **without adopting its pedagogy**

Week 9: Agents and Multi-Step Systems

- Agents as loops and state machines
- Error compounding and termination
- Why autonomy is fragile

Week 10: End-to-End System Design

- Pipelines and responsibility boundaries
- Failure propagation
- Refactoring as design

Week 11: Deployment and Scaling

- Production realism
- Rate limits, backpressure, observability
- Why “it works locally” means nothing

Week 12: Capstone & Professional Framing

- Integrated system build
- Honest evaluation
- Interview-ready explanations
- Defensible claims

Syllabus - Machine Learning with Transparent Abstractions

Everything from above applies, I copy pasted the structure and making changes only relevant to the material

Summary

I am planning this course in (case you want more traditional ML material) in such a way that it trains students to **build, evaluate, and deploy reliable machine learning systems without myth, hype, or fragile intuitions.**

We do **not** want to focus on:

- algorithm catalogs,
- leaderboard chasing,
- or abstract ML mathematics disconnected from real systems.

Instead, we stress:

- **correct mental models,**
- **engineering and evaluation discipline,**
- **system-level thinking,** and

It is not about merely **run models** or **train algorithms**.

While the course uses familiar ML tasks (classification, regression, clustering), the emphasis is always on **behavior, failure modes, and integration**, not on mysticism.

This is not a • “learn ML algorithms in 12 weeks” course

- Kaggle-style competition bootcamp
- math-heavy statistical learning theory class

The course covers (all from AI/LLM course plus)

- deployment and monitoring (optional)

Core Principles

- Metaphor Audits The course actively identifies and dismantles misleading ML metaphors such as:
 - “the model learns concepts”
 - “the model understands patterns”
 - “the algorithm figures it out”
 - “high accuracy proves correctness”

Students should, by the end of the course, be able to:

- detect misleading metaphors
- explain why they are dangerous
- replace them with precise, operational language

This is a **rare but critical professional skill**, even among experienced practitioners.

- Big-Picture Systems Thinking Over Algorithm Worship Students are trained to reason about:

- data pipelines
- preprocessing and feature boundaries
- training vs inference separation
- evaluation loops
- deployment constraints
- monitoring and drift
- failure propagation

ML models are treated correctly as probabilistic components inside engineered systems, not as intelligent entities.

- Just-In-Time Foundations (No Prerequisites) Students can enter with:

- no programming background
- no ML background
- no math background

Foundations (Python, statistics, probability, linear algebra, optimization) are introduced:

- **only when needed**
- **only to the depth required**

What the Syllabus Covers (High-Level)

Weeks 1–2: What Machine Learning Actually Is

- ML as function approximation under uncertainty
- Data and labels as task specification
- Why models do not “understand” data

Week 3: Just-In-Time Python for ML Evidence

- Loading, inspecting, and summarizing data
- Deterministic code vs stochastic ML behavior
- Evidence generation, not “model magic”

Week 4: Supervised Learning Foundations

- Training/validation/test splits
- Loss functions as human-defined objectives
- Training as optimization, not learning concepts

Week 5: Baselines and Simple Models

- Linear models and kNN
- Why baselines are safety checks
- Interpreting errors instead of chasing accuracy

Week 6: Optimization Machinery

- Gradient descent as local numeric improvement
- Hyperparameters as tradeoffs
- Why optimization can succeed and still fail you

Week 7: Generalization and Overfitting

- Why training success lies
- Regularization and early stopping
- Capacity vs data

Week 8: Evaluation and Measurement

- Metrics as estimates, not proofs
- Variance, repeated runs, and slicing
- Why benchmarks mislead

Week 9: Unsupervised Learning and Representation

- Clustering and PCA
- Similarity vs meaning
- Structure creation, not discovery of truth

Week 10: End-to-End ML Pipelines

- Data > features > model > evaluation > decision
- Reproducibility and versioning
- Failure propagation

Week 11: Deployment, Monitoring, and Drift

- Dataset shift and real-world change
- Monitoring signals and alerts
- Why retraining is not a solution by itself

Week 12: Capstone & Professional Framing

- Integrated ML system build
- Honest evaluation and failure analysis
- Interview-ready, defensible explanations